

Beacon Digest -- 2026-09-02

Linux

A note from LWN

Source: *LWN.net*

The online publication industry, as a whole, is struggling, with challenges coming from multiple directions. Thanks to the support of all of you, our readers, LWN would appear to be doing better than most. But the world has changed around us and, in particular, prices have changed considerably. By n...

[\$] A pause for the Python JIT

Source: *LWN.net*

In 2024 the Python 3.13 release added an experimental just-in-time (JIT) compiler to optimize the way that CPython executes Python code. Since then, work has proceeded on the JIT, albeit perhaps less formally than some might like. In June, Python's steering council (SC) put out an announcement that...

Firefox 155 released

Source: *LWN.net*

Version 155 of the Firefox web browser has been released. Notable changes include a count in the address bar of how many ad trackers Firefox has blocked, container reordering, and ensuring that mailto: links are only opened by explicit user actions. There is also a change of the domain used for "cap...

Security updates for Tuesday

Source: *LWN.net*

Security updates have been issued by AlmaLinux (gzip, iperf3, libxml2, mingw-sqlite, mysql:8.4, nginx:1.26, nodejs:24, php, and tar), Debian (expat and libdbd-csv-perl), Fedora (apache-ivy, bind, bluez, bubblewrap, curl, emacs, epiphany, expat, freerdp, gdk-pixbuf2, GitPython, hcloud, kbd, kernel, l...

5 ways to augment security risk management in the AI era

Source: *Red Hat Blog*

IT operations and security teams receive thousands of alerts every day from threat intelligence sources, such as vulnerability scanners, observability tools, Red Hat

Lightspeed, and more. The challenge is identifying, correlating, and addressing impactful alerts from an ongoing, high-volume stream o...

Why the virtualization decision keeps getting deferred

Source: Red Hat Blog

If double-paying during a platform transition is stalling your virtualization migration, removing that first-year cost overlap can get your project moving again. Most infrastructure teams I speak with have spent the best part of 2 years working out what to do about their virtualization platform. The...

How Ask Red Hat earns trust in enterprise AI troubleshooting

Source: Red Hat Blog

A recent review of 562 empirical studies on human trust in AI found that capability, explainability, transparency, and individual user factors consistently predict whether people will rely on an AI-generated answer. In production, the question is narrower and harder: Will I stake my environment on t...

Automate AI red teaming: Large language model risk identification and mitigation

Source: Red Hat Blog

Large language models (LLMs) remain vulnerable to adversarial attacks that bypass safety controls and produce unwanted content. To help engineering and security teams identify these vulnerabilities systematically, Red Hat OpenShift AI includes an automated red teaming (ART) pipeline. This tool conve...

DevOps

Platform engineering maturity: From toolchain to self-service

Source: CNCF Blog

Most platform engineering conversations tend to split into two rooms pretty quickly. The first room is full of teams who don't have a platform yet. Scattered scripts, tribal knowledge, and every team is doing the same...

Observability in Kubernetes: From metrics to meaning

Source: CNCF Blog

Kubernetes made infrastructure more programmable, scalable, and resilient. It also made production systems harder to reason about. Workloads move, replicas churn, dependencies multiply, and a single user request can cross ingress, services, queues, storage, and background...

Kubernetes v1.37: etcd RangeStream Cuts Memory Use on Large List Reads

Source: Kubernetes Blog

I am excited to announce that etcd RangeStream is graduating to beta in Kubernetes v1.37. Paired with etcd v3.7, it reduces the memory the API server and etcd need to read a large collection, and makes peak usage more predictable. The cost of large reads The API server serves most list and watch req...

Secure mainframe access with HashiCorp Boundary

Source: HashiCorp Blog

HashiCorp Boundary modernizes mainframe access with identity-based controls, JIT credentials, and central auditing via a mainframe-adjacent worker.

HashiCorp Vault agentic IAM is now generally available

Source: HashiCorp Blog

HashiCorp Vault Enterprise advances AI agent security with recent enhancements now generally available.

Claude Fable 5.1 watermark: It has a blind spot developers can't ignore

Source: The New Stack

Anthropic launched Claude Fable 5.1 on Tuesday with a statistical signature embedded in its generated text, but developers shouldn't expect The post Claude Fable 5.1 watermark: It has a blind spot developers can't ignore appeared first on The New Stack .

Runway wants to generate software as you use it. Solaris is its first step.

Source: The New Stack

Runway introduced Solaris on Monday, the first model of a new class of AI systems it calls Interface World Models, The post Runway wants to generate software as you use it. Solaris is its first step. appeared first on The New Stack .

Anthropic's Fable 5.1 is a bit cheaper, a bit smarter, and refuses a lot less

Source: The New Stack

On Tuesday, Anthropic launched the latest versions of its flagship Fable and Mythos models. Anthropic promises that the updated models The post Anthropic's

Fable 5.1 is a bit cheaper, a bit smarter, and refuses a lot less appeared first on The New Stack .

Your Mac is now part of Perplexity's AI infrastructure

Source: The New Stack

Perplexity wants its AI agents to use more of the computing power already sitting inside a Mac. The company launched The post Your Mac is now part of Perplexity's AI infrastructure appeared first on The New Stack .

GLM-5.3-Flash vs. GLM-5.3: Time and money, not the spec sheet

Source: The New Stack

I always wonder about the end goal when companies launch products so close together and undercut each other by claiming The post GLM-5.3-Flash vs. GLM-5.3: Time and money, not the spec sheet appeared first on The New Stack .

Meta just beat OpenAI and Google at real-time transcription

Source: The New Stack

Meta's Superintelligence Labs on Tuesday launched Muse Voice Transcribe, a new real-time speech recognition model that, on some benchmarks, outperforms The post Meta just beat OpenAI and Google at real-time transcription appeared first on The New Stack .

Vibe-coded apps are the new shadow IT

Source: The New Stack

Shadow IT used to be a SaaS problem. Someone on the marketing team signed up for a tool, connected it The post Vibe-coded apps are the new shadow IT appeared first on The New Stack .

Meta's Claude Code rival exits beta with three new subscription tiers — and it's pushing hard on price

Source: The New Stack

Meta has formally launched Muse Code out of beta, less than a month after first debuting the coding agent. Alongside The post Meta's Claude Code rival exits beta with three new subscription tiers — and it's pushing hard on price appeared first on The New Stack .

When agents build, deploy, and maintain, persistence becomes the hard problem

Source: The New Stack

Every application has always needed a place to keep its state. That is not new. What is new is who. When agents build, deploy, and maintain, persistence becomes the hard problem appeared first on The New Stack .

Security

Researchers Use Claude to Port Pre-Auth RCE Exploit From One PLC Model to Another

Source: *The Hacker News*

Forescout Research - Vedere Labs said it used Anthropic's Claude to port a working pre-authentication remote code execution (RCE) exploit from one WAGO programmable logic controller (PLC) to another, executing attacker-supplied ARM shellcode on live hardware. The exploit targets CVE-2021-31886, a st...

Attackers Exploit Critical Switchvox Flaw to Deploy Reverse Shells Without Credentials

Source: *The Hacker News*

Threat actors are exploiting a severe security vulnerability in Sangoma Switchvox, an enterprise VoIP platform, that could allow unauthenticated remote code execution. The vulnerability in question is CVE-2026-9586 (CVSS score: 9.3), a critical unauthenticated SQL injection vulnerability in Sangoma...

Authorities Turn Sality's P2P Network Against Itself, Cutting Off New Malware Payloads

Source: *The Hacker News*

The U.S. Department of Justice (DoJ) on Tuesday announced the takedown of a long-standing peer-to-peer (P2P) botnet known as Sality as part of a coordinated law enforcement operation. The effort was undertaken on August 31, 2026, by authorities from the U.S., Bulgaria, Hungary, and Romania, in colla...

Attackers Exploit Critical JFrog Artifactory Flaw to Mint Admin Tokens Days After Disclosure

Source: *The Hacker News*

Threat actors are exploiting a newly patched critical security flaw impacting JFrog Artifactory merely days after public disclosure, according to watchTower. The vulnerability in question is CVE-2026-82329 (CVSS score: 9.8), a case of authentication bypass that could lead to administrative access in...

Breeze Comet Executes Hundreds of Fraudulent Transactions via Brazilian Payment Systems

Source: The Hacker News

Brazilian financial services, retail, and e-commerce organizations have become the target of a financially motivated threat actor dubbed Breeze Comet (formerly UNC5669) since 2024. Google Threat Intelligence Group (GTIG) and Mandiant teams described the threat actor as "specializing in manipulating...

13 Malicious Packagist Packages Target Unpatched iPhones to Steal Crypto Wallet Seeds

Source: The Hacker News

Cybersecurity researchers have identified a set of 13 malicious Composer theme packages on Packagist that are designed to inject JavaScript into Vietnamese movie and comic streaming sites that install those libraries and initiate the deployment of spyware aimed at unpatched iOS devices. "The injecte...

Iranian Hackers Pose as Recruiters to Deliver Cross-Platform RATs Through Coding Tests

Source: The Hacker News

The Iranian Nimbus Manticore hacking group has been attributed to two previously undocumented malware families that highlight the continued evolution of its toolset and likely expand its targeting footprint to infect Linux and Apple macOS systems using cross-platform remote access trojans (RATs) dev...

Threat Actors Don't Want Better Attacks. They Want Repeatable Ones

Source: The Hacker News

The most common way into a company last year was to ask. A web page tells the visitor to prove they are not a robot. While they read the instructions, it quietly places a command on their clipboard. Then it talks them through opening a terminal and pasting it in. The technique is called ClickFix, an...

Attackers Steal METR API Key and Consume AI Credits Worth About \$600,000

Source: The Hacker News

METR (short for Model Evaluation and Threat Research and pronounced "Meter"), a research non-profit that evaluates frontier artificial intelligence (AI) models for their ability to carry out long-horizon, agentic tasks, disclosed that it suffered "two notable security incidents" where external actor...

AI

How AI-native companies turn workflows into operating capability

Source: OpenAI Blog

Basis, Clay, and Exa Labs use AI agents to improve onboarding, account management, and developer integrations. See what enterprise leaders can apply.

Path to Astra: critical capabilities and frontier safeguards

Source: OpenAI Blog

Astra is the first OpenAI model to meet the Critical cybersecurity capability threshold under the Preparedness Framework, with stronger safeguards for release.

Healthcare organizations can now connect EHR and additional industry data to ChatGPT

Source: OpenAI Blog

ChatGPT can now connect to trusted healthcare data, helping clinicians securely access patient context, medical research, and more.

How law firm Gilbert + Tobin governs and scales AI with OpenAI

Source: OpenAI Blog

See how Gilbert + Tobin combines CEO-led commitment, rigorous governance, and human accountability to scale ChatGPT Enterprise and Codex across the firm.

Introducing agentic video understanding with Gemini

Source: Google DeepMind Blog

BenchMIRT: What are LLM benchmarks actually measuring?

Source: Hugging Face Blog

Introducing @huggingface/kernels: 200+ WebGPU Kernels for Local AI

Source: Hugging Face Blog

General Tech

Telli (YC F24) is hiring engineers and designers [Berlin, on-site]

Source: Hacker News (Top 20)

Article URL: <https://careers.telli.com/> Comments URL: <https://news.ycombinator.com/item?id=49533557> Points: 0 # Comments: 0

Commodore 64 released September 1, 1982

Source: Hacker News (Top 20)

Article URL: <https://dfarq.homeip.net/commodore-64-released-september-1-1982/>
Comments URL: <https://news.ycombinator.com/item?id=49533497> Points: 8 # Comments: 0

The Emergent Symbolic Structure of Artificial Neural Networks

Source: Hacker News (Top 20)

Article URL: <https://arxiv.org/abs/2608.29530> Comments URL: <https://news.ycombinator.com/item?id=49531651> Points: 133 # Comments: 41

WebFPGA

Source: Hacker News (Top 20)

Article URL: <https://webfpga.io/> Comments URL: <https://news.ycombinator.com/item?id=49531525> Points: 71 # Comments: 30

True Rate of Unemployment

Source: Hacker News (Top 20)

Article URL: <https://www.lisep.org/tru> Comments URL: <https://news.ycombinator.com/item?id=49530989> Points: 226 # Comments: 181

The efficient frontier of LLM inference

Source: Hacker News (Top 20)

Article URL: <https://www.baseten.co/blog/the-efficient-frontier-of-llm-inference/>
Comments URL: <https://news.ycombinator.com/item?id=49529898> Points: 118 # Comments: 30

FBI Probes Service Selling 153M+ Drivers Licenses

Source: Hacker News (Top 20)

Article URL: <https://krebsonsecurity.com/2026/09/fbi-probes-service-selling-153m-drivers-licenses/> Comments URL: <https://news.ycombinator.com/item?id=49529621> Points: 207 # Comments: 82

My local model setup on an M4 Pro Mac Mini

Source: *Hacker News (Top 20)*

Article URL: <https://lws.io/blog/my-local-model-setup/> Comments URL: <https://news.ycombinator.com/item?id=49529132> Points: 195 # Comments: 106

Show HN: Weedout - Safari extension that hides YouTube AI-labeled videos

Source: *Hacker News (Top 20)*

I'm the developer. I built this because AI-generated conspiracy videos were getting out of control in my own YouTube feed. Weedout is a \$1.99 Safari extension for macOS that removes videos YouTube labels "Made with AI" from feed, search, related videos, playlists and Shorts. It uses YouTube's label...

Show HN: HN Match Maker - Matching "Who Wants to Be Hired?" With "Who's Hiring?"

Source: *Hacker News (Top 20)*

Hi! Every month the "Who Wants to Be Hired?" and "Who's Hiring?" Threads come through I would think to myself, "there should be a match maker for these two threads". So I had abacus.ai whip one up. The methodology is pretty simple, data is extracted from posts using an LLM, score matches based on sa...

Path to Astra: critical capabilities and frontier safeguards

Source: *Hacker News (Top 20)*

Article URL: <https://openai.com/index/path-to-astra/> Comments URL: <https://news.ycombinator.com/item?id=49527595> Points: 155 # Comments: 63

The ChatGPT/Codex app bundles a full copy of LibreOffice

Source: *Hacker News (Top 20)*

Article URL: <https://simonwillison.net/2026/Sep/1/codex-libreoffice/> Comments URL: <https://news.ycombinator.com/item?id=49527396> Points: 408 # Comments: 187

Refurbishing a Tektronix TDS7104 Oscilloscope

Source: *Hacker News (Top 20)*

Article URL: <https://tomverbeure.github.io/2026/08/23/Tektronix-TDS7104-Refurbishing.html> Comments URL: <https://news.ycombinator.com/item?id=49527232> Points: 122 # Comments: 55

How accurate have Ed Zitron's AI skeptic predictions been?

Source: Hacker News (Top 20)

Article URL: <https://danluu.com/zitron/> Comments URL: <https://news.ycombinator.com/item?id=49526069> Points: 692 # Comments: 753

Claude Fable 5.1 and Claude Mythos 5.1

Source: Hacker News (Top 20)

What's new in Claude Fable 5.1 – <https://platform.claude.com/docs/en/models/fable-5-1/whats-n...> System Card: <https://www-cdn.anthropic.com/0339e6a7c5c7b87f5c07798616dc32...> Comments URL: <https://news.ycombinator.com/item?id=49525378> Points: 1232 # Comments: 1148

The creator of Jujutsu has joined ERSC

Source: Hacker News (Top 20)

Article URL: <https://ersc.io/blog/martin-joins-ersc> Comments URL: <https://news.ycombinator.com/item?id=49525297> Points: 232 # Comments: 175

Atlas: A World Model for Spatial Intelligence

Source: Hacker News (Top 20)

Article URL: <https://www.worldlabs.ai/blog/atlas> Comments URL: <https://news.ycombinator.com/item?id=49525160> Points: 222 # Comments: 53

Launch HN: Nori Robotics (YC S26) - A low-cost humanoid robot for development

Source: Hacker News (Top 20)

Hey HN, I'm Antonio from Nori Robotics (<https://norirobotics.com>). We build a \$1,688 bimanual mobile robot in San Francisco for robotics developers and researchers. I started working on Nori while doing robotics research at Columbia. I was teaching robots through human demonstrations, but getting...

Fine, I'll build my own text editor

Source: Hacker News (Top 20)

Article URL: <https://dbushell.com/2026/09/01/text-editor/> Comments URL: <https://news.ycombinator.com/item?id=49524863> Points: 107 # Comments: 95

Show HN: Running 104GB Qwen3.8-Flash-Next on 48GB Mac with at ~12 tok/s

Source: *Hacker News (Top 20)*

I built slotstream, a way to run Qwen3.8-Flash-Next 4-bit on a low-memory mac starting from 16GB, a 125B parameter model that would need 100GB+ memory/RAM, thanks to expert-offloading/ssd-streaming. Easy to install/update, and mac-native using MLX and Swift. It ships with auto-mode, which makes a go...